

Pet Classification with CNNs


Project Overview

The goal of this project is to implement a neural network that classifies images of 37 breeds of cats and dogs from the Oxford-IIIT-Pet dataset. The project is divided into two tasks: first, you will be asked to implement from scratch your own neural network for image classification; then, you will fine-tune a pre-trained network provided by Keras.

Dataset

Overview

The dataset you will be using contains images of pets.



You can download the dataset at the following URL: <https://www.robots.ox.ac.uk/~vgg/data/pets/>.

Content

The repository contains a 37-category pet dataset with roughly 200 images for each class. The images have a large variation in scale, pose and lighting. All images have an associated ground truth annotation of breed. The dataset is already split into training, validation and testing. You should use only the training and validation splits for development in both the first and second tasks.

First Task: design your own network

The goal is to implement a convolutional neural network for image classification and train it on the Oxford-IIIT-Pet Dataset. You should consider yourself satisfied once you obtain a classification accuracy on the validation split of around 60%. You are free to achieve that however you want, except for a few rules you must follow:

- You cannot simply instantiate an off-the-shelf Keras network. Instead, you must construct your network as a composition of existing Keras layers;
- Justify every design choice you make. Design choices include network architecture, training hyperparameters, and, possibly, dataset preprocessing steps. You can either (i) start from the simplest convolutional network you can think of and add complexity one step at a time, while showing how each step gets you closer to the target 60%, or (ii) start from a model that is already able to achieve the desired accuracy and show how, by removing some of its components, its performance changes. You can show your results however you want: training plots, console-printed values or tables.

Don't be too concerned with your network performance: the 60% is just to give you an idea of when to stop.

Second Task: fine-tune an existing network

Your goal is to fine-tune a pre-trained ResNet-18 model on the Oxford-IIIT-Pet Dataset. Use the implementation provided by Keras; do not implement it yourselves (i.e. exactly what you could not do in the first task). Specifically, you must use the Keras ResNet-18 model pre-trained on ImageNet-1K.

Divide your fine-tuning into two parts:

1. First, fine-tune the ResNet-18 with the same training hyperparameters you used for your best model in the first task of the project;
2. Then, tweak the training hyperparameters in order to increase the accuracy on the validation split of the Oxford-IIIT-Pet Dataset. Justify your choices by analysing the training plots and/or citing sources that guided you in your decisions (papers, blog posts, YouTube videos, or whatever else you find enlightening). You should consider yourself satisfied once you obtain a classification accuracy on the validation split between 80 and 90%.

Final Task: assess the performance on the testing set

Assess the performance of your best model from both tasks on the testing set as the final evaluation.